

EECS 111:

System Software

Lecture : CPU Scheduling

Prof. Mohammad Al Faruque

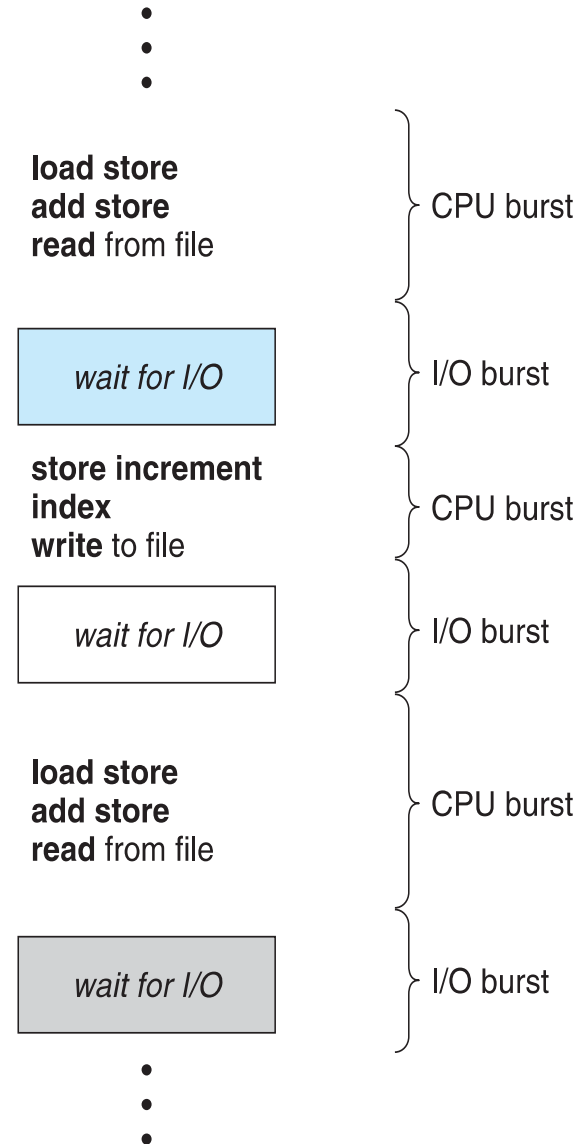
**The Henry Samueli School of Engineering
Electrical Engineering & Computer Science
University of California Irvine (UCI)**

Chapter 6: CPU Scheduling

- ❑ **Basic Concepts**
- ❑ **Scheduling Criteria**
- ❑ **Scheduling Algorithms**
- ❑ **Real-Time CPU Scheduling**

Basic Concepts

- ❑ Maximum CPU utilization obtained with **multiprogramming**
- ❑ CPU-I/O Burst Cycle – Process execution consists of a **cycle** of CPU execution and I/O wait
- ❑ **CPU burst** followed by **I/O burst**
- ❑ CPU burst distribution is of main concern



Histogram of CPU-burst Times


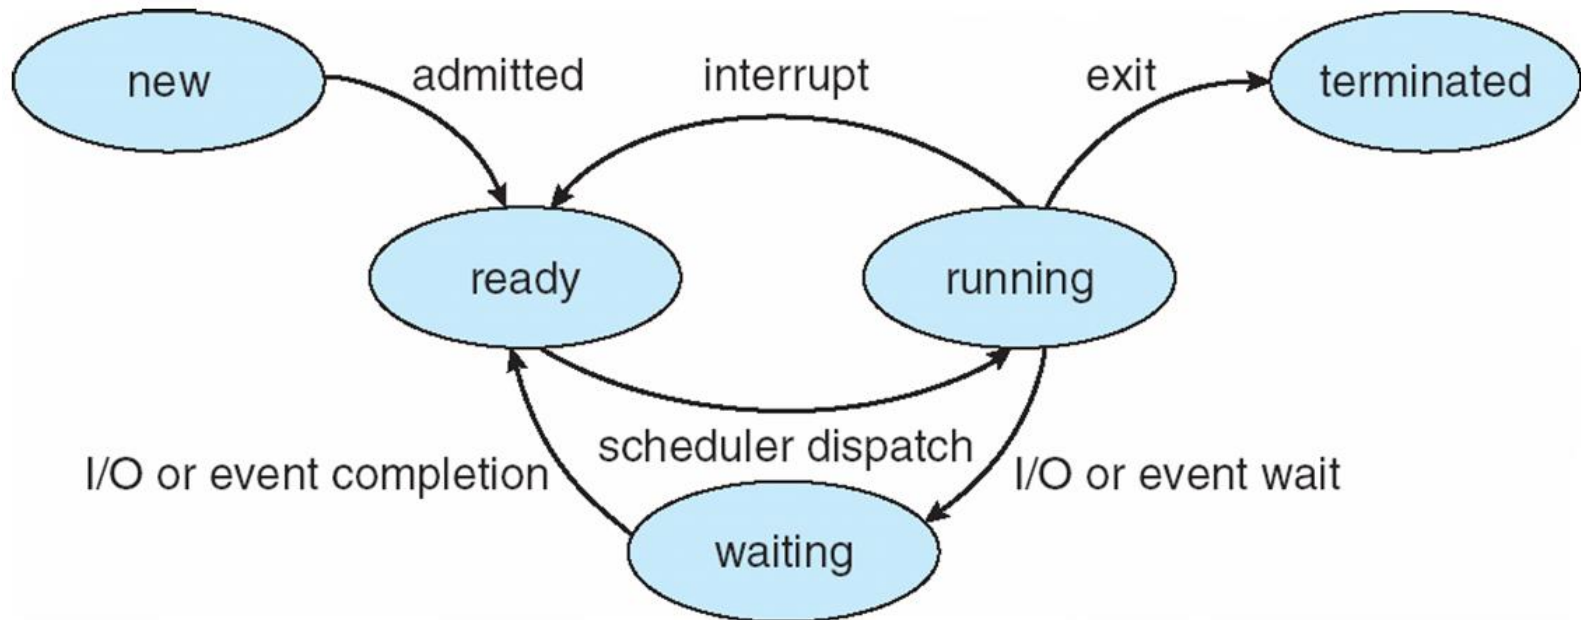


Diagram of Process State



Schedulers

❑ Long-term scheduler (or job scheduler) -

- ❑ selects which processes should be brought into the ready queue.
- ❑ invoked very infrequently (seconds, minutes); may be slow.
- ❑ controls the degree of multiprogramming

❑ Short term scheduler (or CPU scheduler) -

- ❑ selects which process should execute next and allocates CPU.
- ❑ invoked very frequently (milliseconds) - must be very fast

❑ Medium Term Scheduler

- ❑ swaps out process temporarily → swapping
- ❑ balances load for better throughput

CPU Scheduler

- ❑ **Short-term scheduler** selects from among the processes in ready queue, and allocates the CPU to one of them
 - ❑ Queue may be ordered in various ways
- ❑ CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- ❑ **Scheduling under 1 and 4 is nonpreemptive**
- ❑ All other scheduling is **preemptive**
 - ❑ Consider access to shared data
 - ❑ Consider preemption while in kernel mode
 - ❑ Consider interrupts occurring during crucial OS activities

Dispatcher

- ❑ **Dispatcher module** gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - ❑ switching context
 - ❑ switching to user mode
 - ❑ jumping to the proper location in the user program to restart that program

- ❑ **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running

Scheduling Criteria

☐ CPU utilization

- ☐ keep the CPU as busy as possible

☐ Throughput

- ☐ # of processes that complete their execution per time unit

☐ Turnaround time

- ☐ amount of time to execute a particular process

☐ Waiting time

- ☐ amount of time a process has been waiting in the ready queue

☐ Response time

- ☐ amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

First-Come, First-Served (FCFS) Scheduling

10

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- ❑ Suppose that the processes arrive in the order: P_1 , P_2 , P_3
The Gantt Chart for the schedule is:



- ❑ **Waiting time**

❑ for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$

- ❑ **Average waiting time: $(0 + 24 + 27)/3 = 17$**

FCFS Scheduling (Cont.)

- ❑ Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

- ❑ The Gantt chart for the schedule is:



- ❑ Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- ❑ **Average waiting time:** $(6 + 0 + 3)/3 = 3$
- ❑ Much better than previous case
- ❑ Convoy effect - short process behind long process
 - ❑ Consider one CPU-bound and many I/O-bound processes

Shortest-Job-First (SJF) Scheduling


- ❑ Associate with each process the length of its next CPU burst
 - ❑ Use these lengths to schedule the process with the shortest time

- ❑ **SJF is optimal** – gives minimum average waiting time for a given set of processes
 - ❑ The difficulty is knowing the length of the next CPU request
 - ❑ Could ask the user

Example of SJF

<u>Process</u>	<u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

❑ SJF scheduling chart



❑ Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$

Determining Length of Next CPU Burst

- ❑ Can only estimate the length – should be similar to the previous one
 - ❑ Then pick process with shortest predicted next CPU burst
- ❑ Can be done by using the length of previous CPU bursts, using **exponential averaging**


1. t_n = actual length of n^{th} CPU burst
2. τ_{n+1} = predicted value for the next CPU burst
3. $\alpha, 0 \leq \alpha \leq 1$
4. Define :

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n.$$

- ❑ Commonly, α set to $\frac{1}{2}$
- ❑ Preemptive version called **shortest-remaining-time-first**

Prediction of the Length of the Next CPU Burst

15




CPU burst (t_i)	6	4	6	4	13	13	13	...
"guess" (τ_i)	10	8	6	5	9	11	12	...

Example of Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

- Preemptive SJF Gantt Chart**



- Average waiting time = $[(10-1)+(1-1)+(17-2)+5-3]/4 = 26/4 = 6.5$ msec**


Priority Scheduling

- ❑ A priority number (integer) is associated with each process
- ❑ The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - ❑ Preemptive
 - ❑ Nonpreemptive
- ❑ SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- ❑ Problem \equiv **Starvation** – low priority processes may never execute
- ❑ Solution \equiv **Aging** – as time progresses increase the priority of the process

Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

□ Priority scheduling Gantt Chart



□ Average waiting time = 8.2 msec

Round Robin (RR)

- ❑ Each process gets a small unit of CPU time (**time quantum q**), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- ❑ If there are **n processes** in the ready queue and the **time quantum is q** , then each **process gets $1/n$ of the CPU time** in chunks of at most q time units at once.
- ❑ **No process waits more than $(n-1)q$ time units.**
- ❑ Timer interrupts every quantum to schedule next process
- ❑ Performance
 - ❑ q large \Rightarrow FIFO
 - ❑ q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high

Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3


❑ The Gantt chart is:



- ❑ Typically, higher average turnaround than SJF, but better *response*
- ❑ **q should be large compared to context switch time**
- ❑ q usually 10ms to 100ms, context switch < 10 usec

Time Quantum and Context Switch Time

21



Scheduling Criteria

☐ CPU utilization

- ☐ keep the CPU as busy as possible

☐ Throughput

- ☐ # of processes that complete their execution per time unit

☐ Turnaround time

- ☐ amount of time to execute a particular process


☐ Waiting time

- ☐ amount of time a process has been waiting in the ready queue

☐ Response time

- ☐ amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

Turnaround Time Varies With The Time Quantum

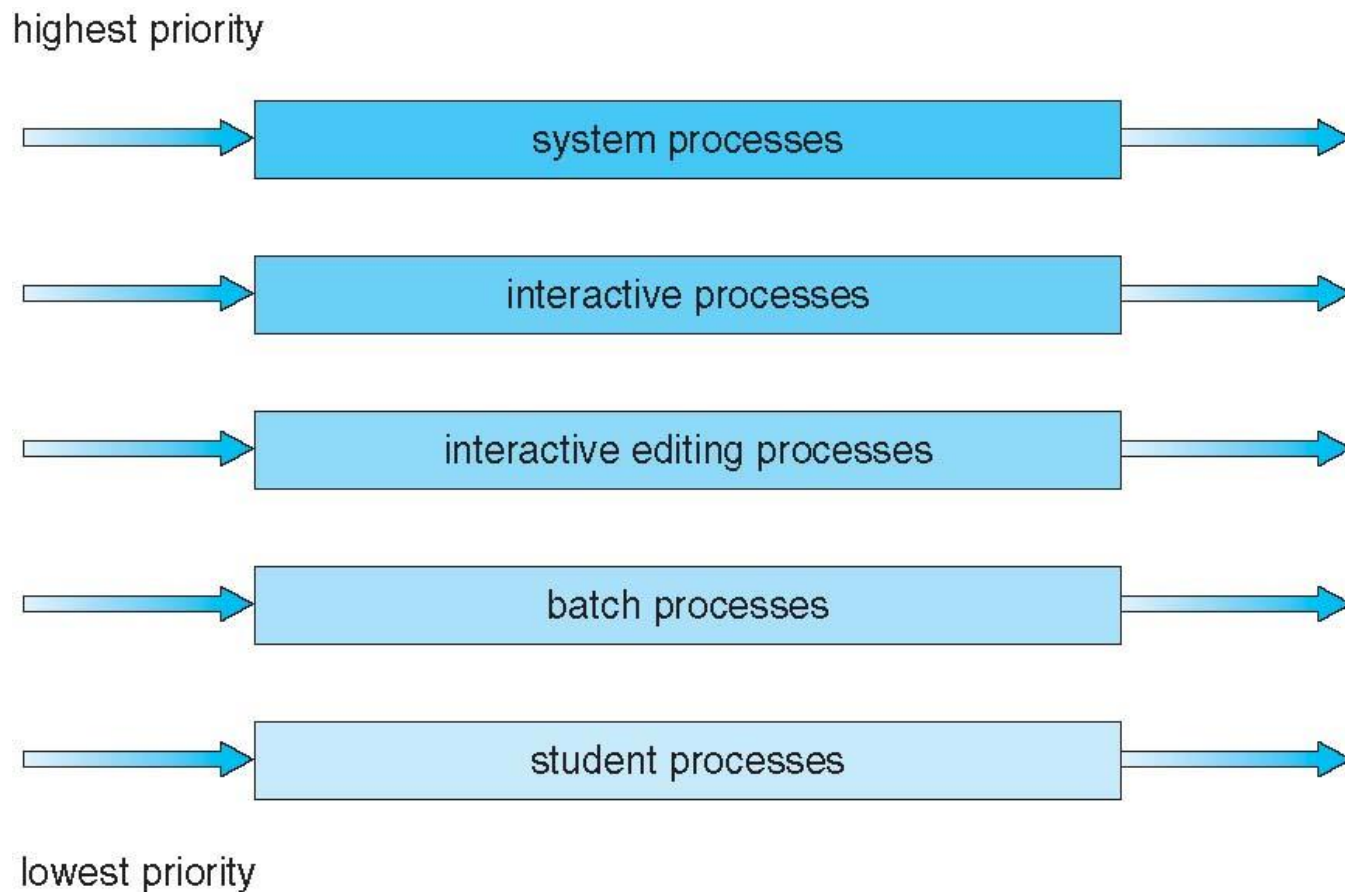


process	time
P_1	6
P_2	3
P_3	1
P_4	7

Multilevel Queue

- ❑ Ready queue is partitioned into separate queues, e.g.:
 - ❑ **foreground** (interactive)
 - ❑ **background** (batch)
- ❑ Process permanently in a given queue
- ❑ Each queue has its own scheduling algorithm:
 - ❑ foreground – RR
 - ❑ background – FCFS
- ❑ Scheduling must be done between the queues:
 - ❑ **Fixed priority scheduling**; (i.e., serve all from foreground then from background) → **Possibility of starvation.**
 - ❑ **Time slice** – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e.,
 - ❑ 80% to foreground in RR
 - ❑ 20% to background in FCFS

Multilevel Queue Scheduling



Multilevel Feedback Queue

- ❑ A process can move between the various queues → aging can be implemented this way

- ❑ Multilevel-feedback-queue scheduler defined by the following parameters:
 - ❑ number of queues
 - ❑ scheduling algorithms for each queue
 - ❑ method used to determine when to upgrade a process
 - ❑ method used to determine when to demote a process
 - ❑ method used to determine which queue a process will enter when that process needs service


Example of Multilevel Feedback Queue

Three queues:

- ❑ Q_0 – RR with time quantum 8 milliseconds
- ❑ Q_1 – RR time quantum 16 milliseconds
- ❑ Q_2 – FCFS

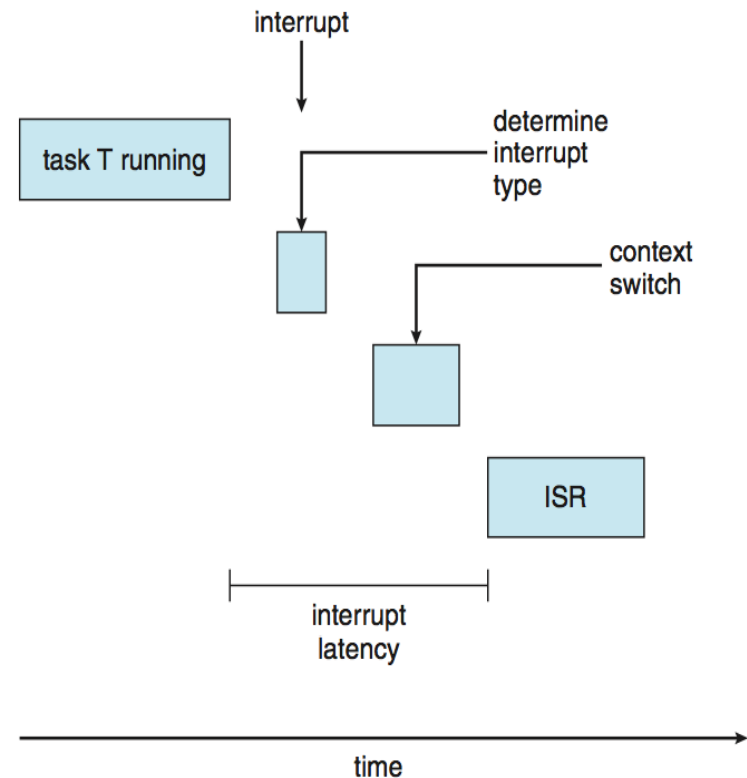
Scheduling

- ❑ A new job enters queue Q_0 which is served RR
 - ❑ When it gains CPU, job receives 8 milliseconds
 - ❑ If it does not finish in 8 milliseconds, job is moved to queue Q_1
- ❑ At Q_1 job is again served RR and receives 16 additional milliseconds
 - ❑ If it still does not complete, it is preempted and moved to queue Q_2
- ❑ At Q_2 job is served FCFS only if Q_0 and Q_1 are empty



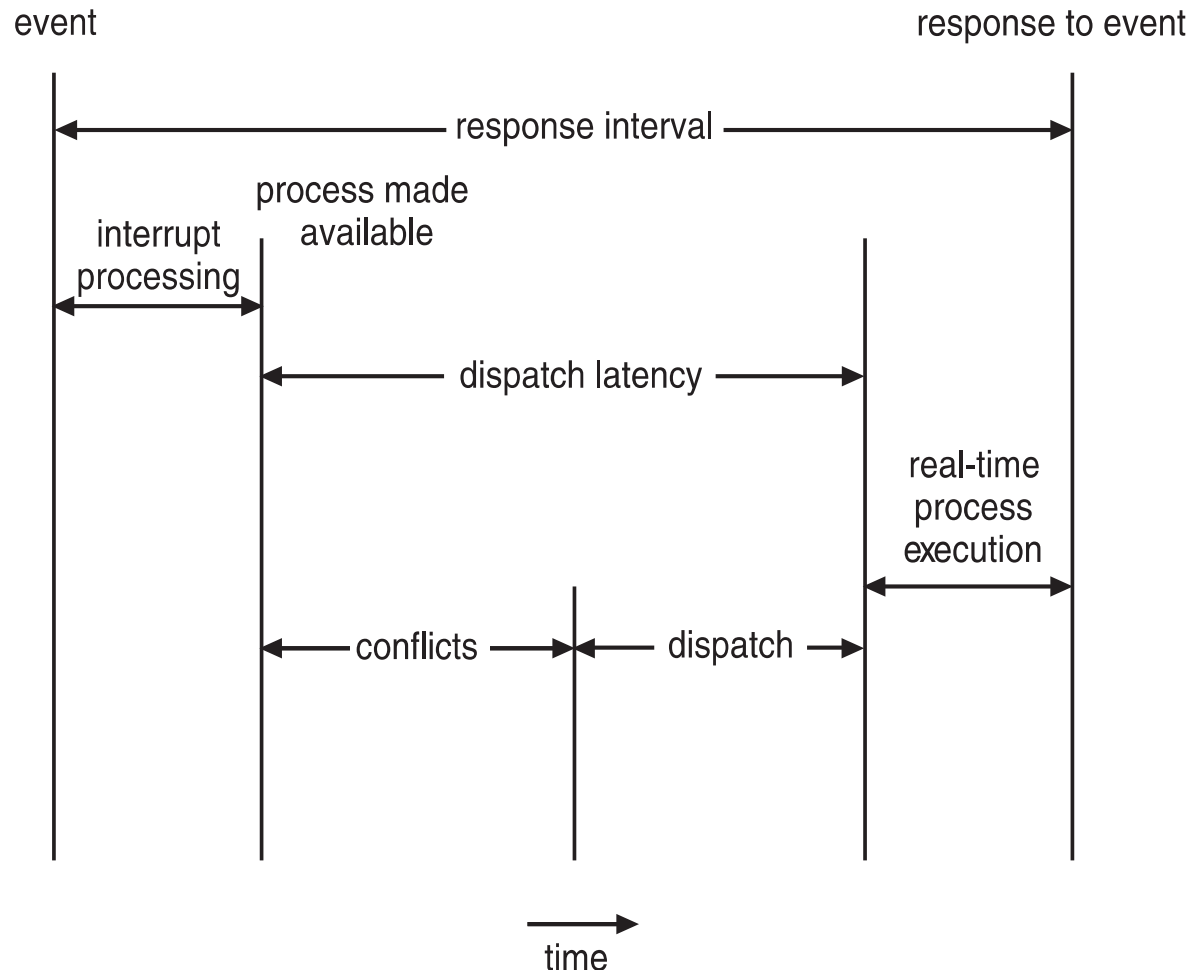
Real-Time CPU Scheduling

- ❑ Can present obvious challenges
- ❑ **Soft real-time systems** – no guarantee as to when critical real-time process will be scheduled
- ❑ **Hard real-time systems** – task must be serviced by its deadline
- ❑ Two types of latencies affect performance
 1. **Interrupt latency** – time from arrival of interrupt to start of routine that services interrupt
 2. **Dispatch latency** – time for schedule to take current process off CPU and switch to another




Real-Time CPU Scheduling (Cont.)

- ❑ **Conflict phase of dispatch latency:**
 1. Preemption of any process running in kernel mode
 2. Release by low-priority process of resources needed by high-priority processes



Priority-based Scheduling

- ❑ For real-time scheduling, scheduler must support preemptive, priority-based scheduling
 - ❑ But only guarantees soft real-time
- ❑ For hard real-time must also provide ability to meet deadlines
- ❑ Processes have new characteristics: **periodic** ones require CPU at constant intervals
 - ❑ Has processing time t , deadline d , period p
 - ❑ $0 \leq t \leq d \leq p$
 - ❑ **Rate** of periodic task is $1/p$



Rate Monotonic Scheduling

- ❑ A priority is assigned based on the inverse of its period
- ❑ Shorter periods = higher priority;
- ❑ Longer periods = lower priority
- ❑ P_1 is assigned a higher priority than P_2 .

$P_1 = 50$


$t_1 = 20$

$P_2 = 100$


$t_2 = 35$

Deadline for each process requires that it complete its CPU burst by the start of its next period.

Start of the period of P_1



Missed Deadlines with Rate Monotonic Scheduling



Start of the period of P_1

$$P_1 = 50$$

$$t_1 = 25$$

$$P_2 = 80$$

$$t_2 = 35$$

Earliest Deadline First Scheduling (EDF)

□ Priorities are assigned according to deadlines:


the earlier the deadline, the higher the priority;
the later the deadline, the lower the priority

$$P_1 = 50$$

$$t_1 = 25$$

$$P_2 = 80$$

$$t_2 = 35$$



References

Part of the contents of this lecture has been adapted from the book Abraham Silberschatz, Peter B. Galvin, Greg Gagne: "Operating System Concept ", Publisher : Wiley; 9 edition (December 17, 2012), ISBN-13: 978-1118063330

Slides also contain lecture materials from John Kubiawicz (Berkeley), John Ousterhout (Stanford), Nalini (UCI), Rainer (UCI), and others

Some slides adapted from <http://www-inst.eecs.berkeley.edu/~cs162/> Copyright © 2010 UCB

**Thank you for your
attention**