

# **EECS 111:**

## **System Software**

### **Lecture : File System Implementation**

**Prof. Mohammad Al Faruque**

**The Henry Samueli School of Engineering  
Electrical Engineering & Computer Science  
University of California Irvine (UCI)**

# File System Implementation

- ❑ **File-System Structure**
- ❑ **File-System Implementation**
- ❑ **Directory Implementation**
- ❑ **Allocation Methods**

# File-System Structure

- ❑ **File structure**
  - ❑ Logical storage unit
  - ❑ Collection of related information
- ❑ **File system** resides on secondary storage (disks)
  - ❑ Provided user interface to storage, mapping logical to physical
  - ❑ Provides efficient and convenient access to disk by allowing data to be stored, located retrieved easily
- ❑ **Disk** provides in-place rewrite and random access
  - ❑ I/O transfers performed in **blocks** of **sectors** (usually 512 bytes)
- ❑ **File control block** – storage structure consisting of information about a file
- ❑ **Device driver** controls the physical device

**File system organized into layers**

# File-System Structure

- ❑ **File structure**

- ❑ Logical storage unit
- ❑ Collection of related information

- ❑ **File system** resides on secondary storage (disks)

- ❑ Provided user interface to storage, mapping logical to physical


## Challenges of a file system

- ❑ How the file should look to the user? → previous lecture
- ❑ Creating algorithms and data structures to map the logical file system onto the physical secondary-storage device?

- ❑ **Device driver** controls the physical device

**File system organized into layers**

# Layered File System



# File System Layers

- ❑ **Device drivers** manage I/O devices at the I/O control layer
  - ❑ Given commands like “**read drive1, cylinder 72, track 2, sector 10, into memory location 1060**” outputs low-level hardware specific commands to hardware controller
- ❑ **Basic file system** given command like “**retrieve block 123**” translates to device driver
  - ❑ Also manages memory buffers and caches (allocation, freeing, replacement)
  - ❑ Buffers hold data in transit
  - ❑ Caches hold frequently used data
- ❑ **File organization module** understands files, logical address, and physical blocks
  - ❑ Translates logical block # to physical block #
  - ❑ Manages free space, disk allocation

# File System Layers (Cont.)

- ❑ **Logical file system manages metadata information**
  - ❑ Translates file name into file number, file handle, location by maintaining file control blocks (**inodes** in Unix)
  - ❑ Directory management
  - ❑ Protection
- ❑ **Layering useful for reducing complexity and redundancy, but adds overhead and can decrease performance**
  - ❑ Logical layers can be implemented by any coding method according to OS designer
- ❑ **Many file systems, sometimes many within an operating system**
  - ❑ Each with its own format (CD-ROM is ISO 9660; Unix has **UFS**, FFS; Windows has FAT, FAT32, NTFS as well as floppy, CD, DVD Blu-ray, Linux has more than 40 types, with **extended file system** ext2 and ext3 leading; plus distributed file systems, etc.)
  - ❑ New ones still arriving – ZFS, GoogleFS, Oracle ASM, FUSE

# File System Implementation

- ❑ File-System Structure
- ❑ **File-System Implementation**
- ❑ Directory Implementation
- ❑ Allocation Methods
- ❑ Free-Space Management



# File-System Implementation

- ❑ We have system calls at the API level, but how do we implement their functions?
  - ❑ **On-disk and in-memory structures**
    - On-disk structures
- ❑ **Boot control block** contains info needed by system to boot OS from that volume
  - ❑ Needed if volume contains OS, usually first block of volume
- ❑ **Volume control block (superblock, master file table)** contains volume details
  - ❑ Total # of blocks, # of free blocks, block size, free block pointers or array
- ❑ Directory structure organizes the files
  - ❑ Names and inode numbers, master file table
- ❑ Per-file **File Control Block (FCB)** contains many details about the file
  - ❑ Inode number, permissions, size, dates
  - ❑ NFTS stores into in master file table using relational DB structures

# A Typical File Control Block

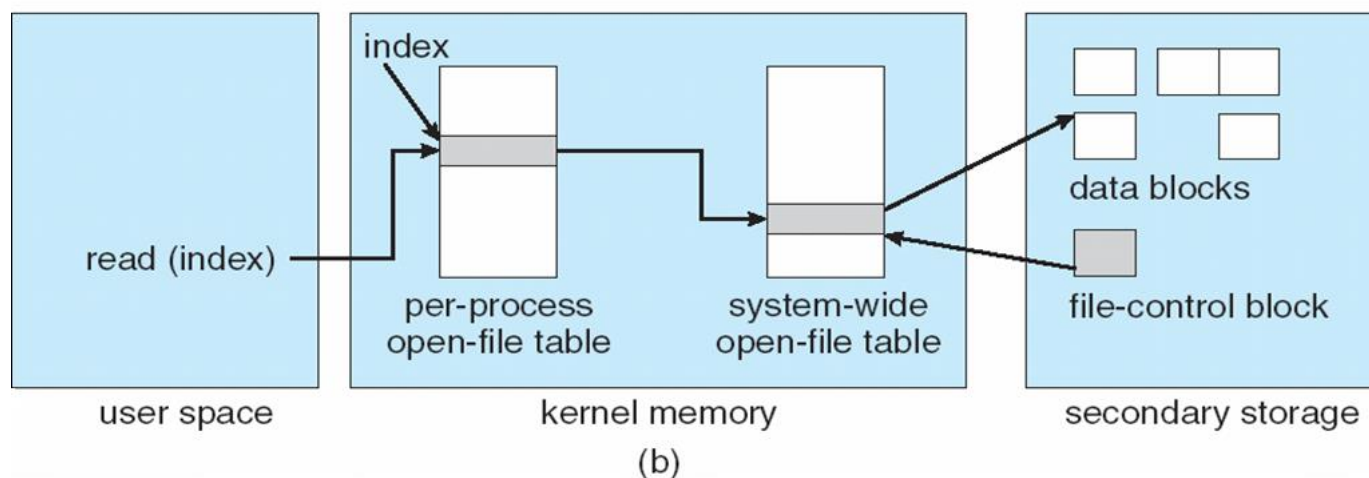

|  |
|--|
| file permissions                                 |
| file dates (create, access, write)               |
| file owner, group, ACL                           |
| file size  |
| file data blocks or pointers to file data blocks |

# In-Memory File System Structures

## In-memory structures

- ❑ Mount table storing file system mounts, mount points, file system types
- ❑ An in-memory directory-structure cache holds the directory information of recently accessed directories.
- ❑ System-wide open file table contains a copy of the FCB of each open file, also other information
- ❑ Per-process Open-file Table pointer to the appropriate entry in the system-wide open file table, other information
- ❑ Buffers hold file-system blocks when are being read from disk or write to disk.

# In-Memory File System Structures



# File System Implementation

- ❑ File-System Structure
- ❑ File-System Implementation
- ❑ Directory Implementation
- ❑ **Allocation Methods**
- ❑ Free-Space Management

# Allocation Methods

- ❑ **An allocation method refers to how disk blocks are allocated for files:**
  
- ❑ **Challenge:** The major challenge is how to allocate space to these files so that disk space is utilized effectively and files can be accessed quickly?
  
- ❑ **Three methods exist:**
  1. Contiguous
  2. Linked
  3. Indexed

# Allocation Methods - Contiguous

- ❑ **Contiguous allocation** – each file occupies set of contiguous blocks
  - ❑ Best performance in most cases
  - ❑ Simple – only starting location (block #) and length (number of blocks) are required
  - ❑ Problems include finding space for file, knowing file size, external fragmentation, need for **compaction off-line (downtime)** or **on-line**
  - ❑ **Second problem** → how much space is needed for a file?
  - ❑ **Similar to problems seen in chapter 8 memory allocation** → first fit, best fit, worst fit

# Contiguous Allocation

## ❑ Mapping from logical to physical

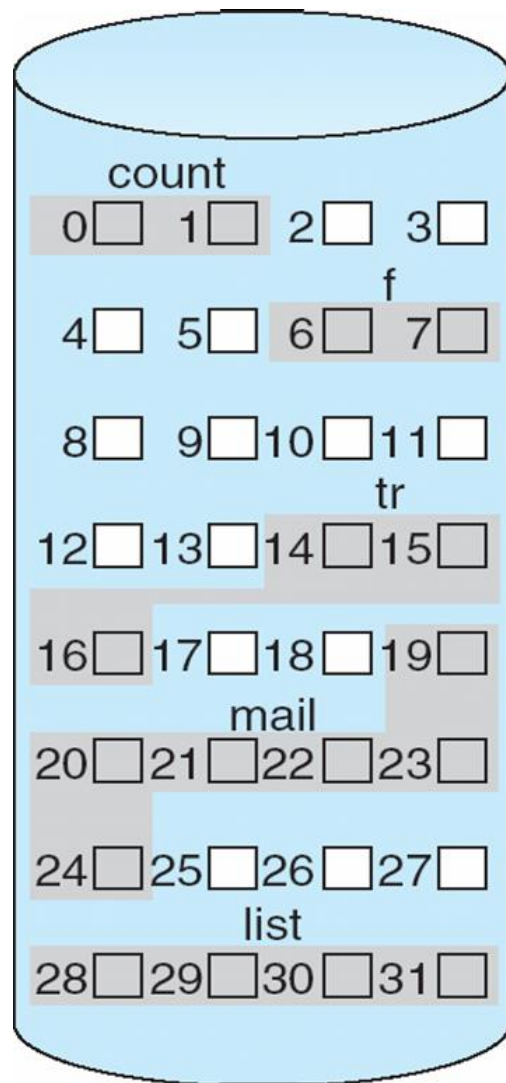


Block to be accessed =  $Q + \text{starting address}$

Displacement into block =  $R$



# Contiguous Allocation of Disk Space



directory

| file  | start | length |
|-------|-------|--------|
| count | 0     | 2      |
| tr    | 14    | 3      |
| mail  | 19    | 6      |
| list  | 28    | 4      |
| f     | 6     | 2      |

# Extent-Based Systems


- ❑ Many newer file systems (i.e., Veritas File System) use a modified contiguous allocation scheme
- ❑ Extent-based file systems allocate disk blocks in extents
- ❑ An **extent** is a contiguous block of disks
  - ❑ Extents are allocated for file allocation
  - ❑ A file consists of one or more extents

# Allocation Methods - Linked

- ❑ **Linked allocation** – each file a linked list of blocks
  - ❑ File ends at **nil** pointer
  - ❑ No external fragmentation
  - ❑ Each block contains pointer to next block
  - ❑ No compaction, external fragmentation
  - ❑ Free space management system called when new block needed
  - ❑ Improve efficiency by clustering blocks into groups but increases internal fragmentation
  - ❑ Reliability can be a problem
  - ❑ Locating a block can take many I/Os and disk seeks


# Linked Allocation

- ❑ Each file is a linked list of disk blocks: blocks may be scattered anywhere on the disk



# Linked Allocation


## ❑ Mapping



Block to be accessed is the  $Q^{\text{th}}$  block in the linked chain of blocks representing the file.

Displacement into block =  $R + 1$


# Linked Allocation



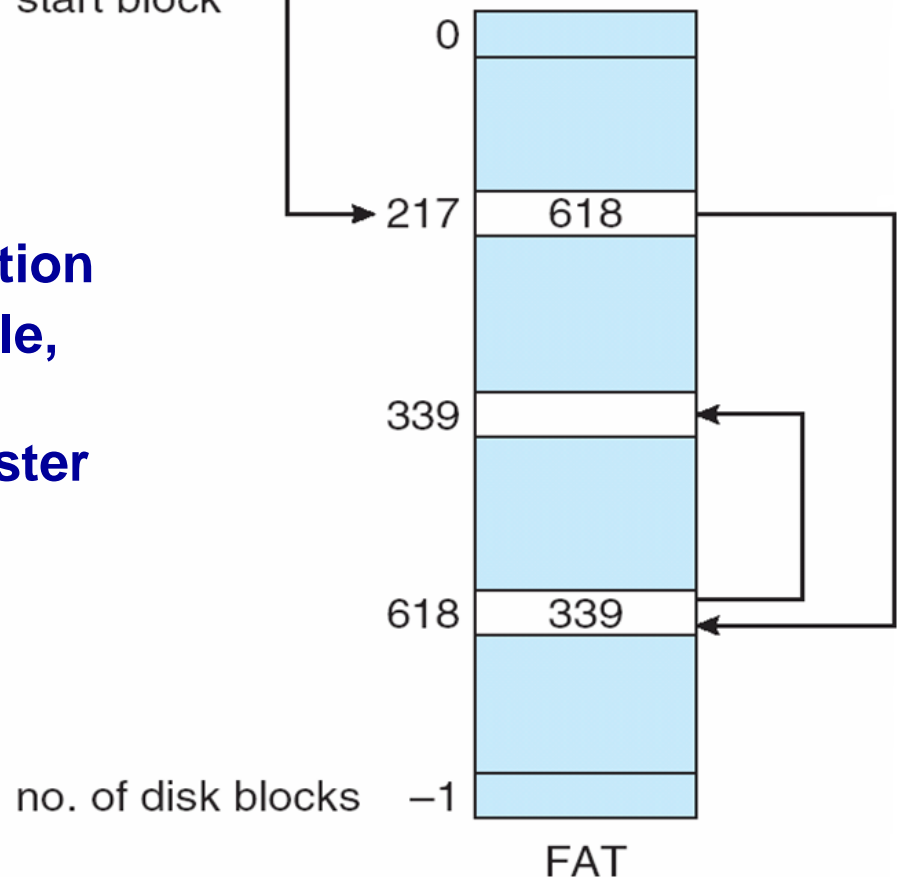
❑ What is the problem in linked allocation? Study!

# File-Allocation Table

directory entry



- ❑ **FAT (File Allocation Table) variation**
  - ❑ Beginning of volume has table, indexed by block number
  - ❑ Much like a linked list, but faster on disk and cacheable
  - ❑ New block allocation simple
  - ❑ **What is the problem in FAT?**  
**Study!**




# Allocation Methods - Indexed

## ❑ Indexed allocation

- ❑ Each file has its own **index block(s)** of pointers to its data blocks


## ❑ Logical view



index table



# Example of Indexed Allocation



# References

Part of the contents of this lecture has been adapted from the book Abraham Silberschatz, Peter B. Galvin, Greg Gagne: "Operating System Concept ", Publisher : Wiley; 9 edition (December 17, 2012), ISBN-13: 978-1118063330

Slides also contain lecture materials from John Kubiawicz (Berkeley), John Ousterhout (Stanford), Nalini (UCI), Rainer (UCI), and others

Some slides adapted from <http://www-inst.eecs.berkeley.edu/~cs162/> Copyright © 2010 UCB

**Thank you for your  
attention**